

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv('/content/drive/MyDrive/Colab Notebooks/diabetes.csv')
```

df

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
0	6	148	72	35	0	33.6		0.627	50	1
1	1	85	66	29	0	26.6		0.351	31	0
2	8	183	64	0	0	23.3		0.672	32	1
3	1	89	66	23	94	28.1		0.167	21	0
4	0	137	40	35	168	43.1		2.288	33	1
...
763	10	101	76	48	180	32.9		0.171	63	0
764	2	122	70	27	0	36.8		0.340	27	0
765	5	121	72	23	112	26.2		0.245	30	0
766	1	126	60	0	0	30.1		0.349	47	1
767	1	93	70	31	0	30.4		0.315	23	0

768 rows × 9 columns

df.shape

(768, 9)

df.describe()

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000		768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578		0.471876	33.240885	0.3485
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160		0.331329	11.760232	0.4765
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000		0.078000	21.000000	0.0000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000		0.243750	24.000000	0.0000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000		0.372500	29.000000	0.0000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000		0.626250	41.000000	1.0000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000		2.420000	81.000000	1.0000

```
missing_cols = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
```

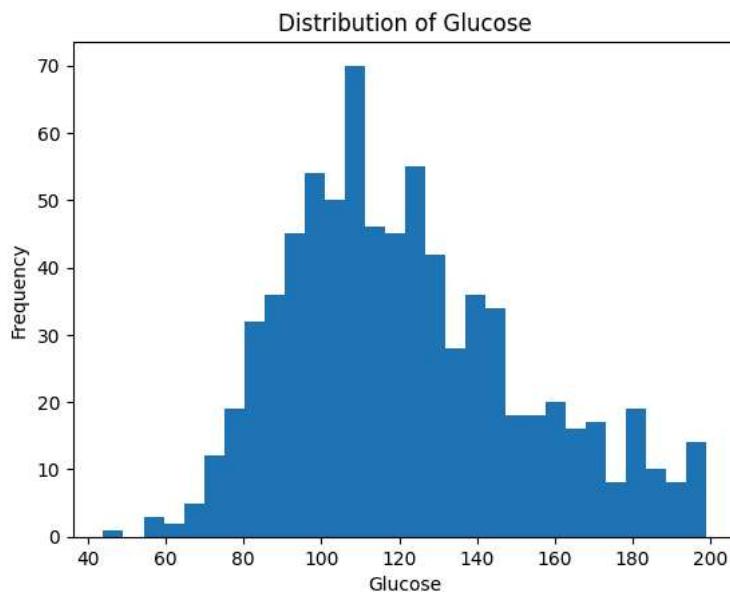
```
for col in missing_cols:
    df[col] = df[col].replace(0, np.nan) # Replace 0 with np.nan (null) in each specified column
```

df.isnull().sum()

```
          0  
Pregnancies      0  
Glucose          5  
BloodPressure    35  
SkinThickness   227  
Insulin         374  
BMI             11  
DiabetesPedigreeFunction  0  
Age              0  
Outcome         0
```

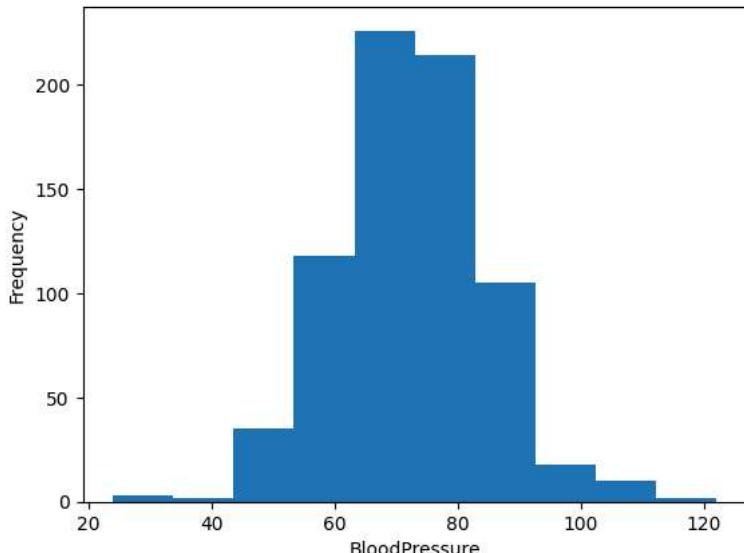
dtype: int64

```
plt.hist(df['Glucose'].dropna(), bins=30)  
plt.title(f"Distribution of Glucose")  
plt.xlabel("Glucose")  
plt.ylabel("Frequency")  
plt.grid(False)  
plt.show()
```



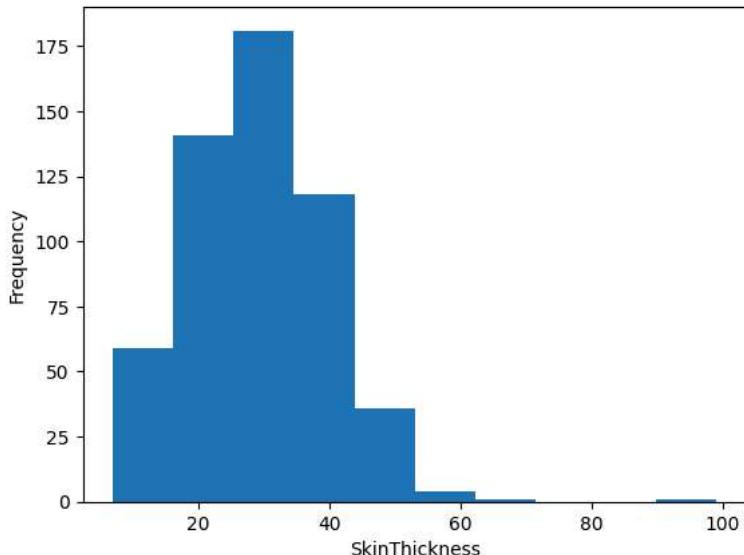
```
plt.hist(df['BloodPressure'].dropna(), bins=10)  
plt.title(f"Distribution of BloodPressure")  
plt.xlabel("BloodPressure")  
plt.ylabel("Frequency")  
plt.grid(False)  
plt.show()
```

Distribution of BloodPressure



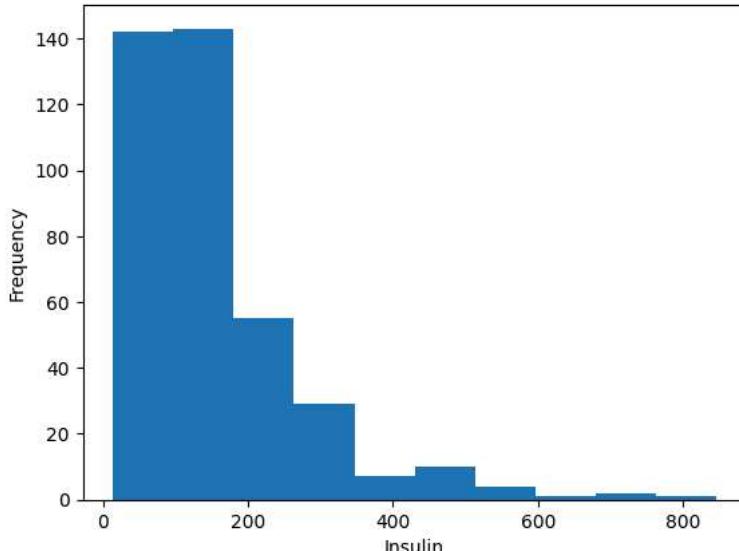
```
plt.hist(df['SkinThickness'].dropna(), bins=10)
plt.title(f"Distribution of SkinThickness")
plt.xlabel("SkinThickness")
plt.ylabel("Frequency")
plt.grid(False)
plt.show()
```

Distribution of SkinThickness



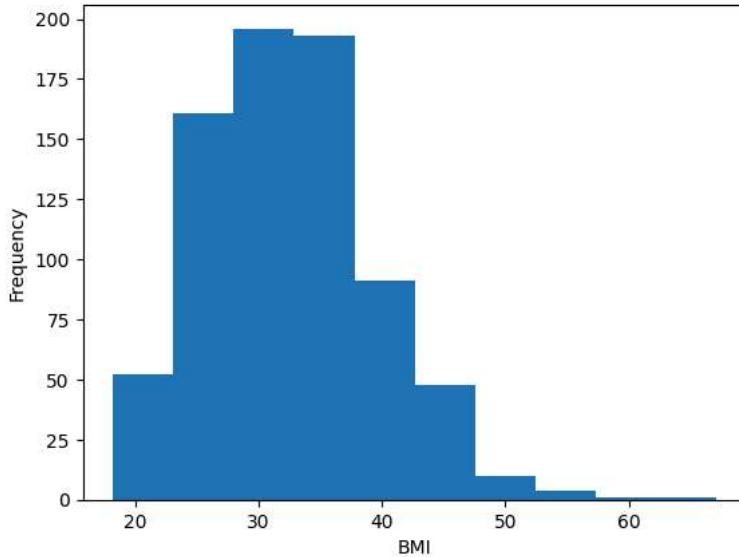
```
plt.hist(df['Insulin'].dropna(), bins=10)
plt.title(f"Distribution of Insulin")
plt.xlabel("Insulin")
plt.ylabel("Frequency")
plt.grid(False)
plt.show()
```

Distribution of Insulin



```
plt.hist(df['BMI'].dropna(), bins=10)
plt.title(f"Distribution of BMI")
plt.xlabel("BMI")
plt.ylabel("Frequency")
plt.grid(False)
plt.show()
```

Distribution of BMI



```
df['Glucose'] = df['Glucose'].fillna(df['Glucose'].mean()) # Replace missing values in the 'Glucose' column

df['BloodPressure'] = df['BloodPressure'].fillna(df['BloodPressure'].mean()) # Replace missing values in the 'BloodPressure'

df['SkinThickness'] = df['SkinThickness'].fillna(df['SkinThickness'].median()) # Replace missing values in the 'SkinThickness'

df['Insulin'] = df['Insulin'].fillna(df['Insulin'].median()) # Replace missing values in the 'Insulin' column

df['BMI'] = df['BMI'].fillna(df['BMI'].mean()) # Replace missing values in the 'BMI' column with mean

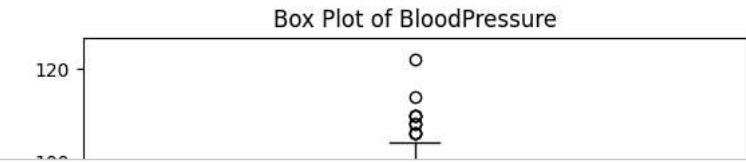
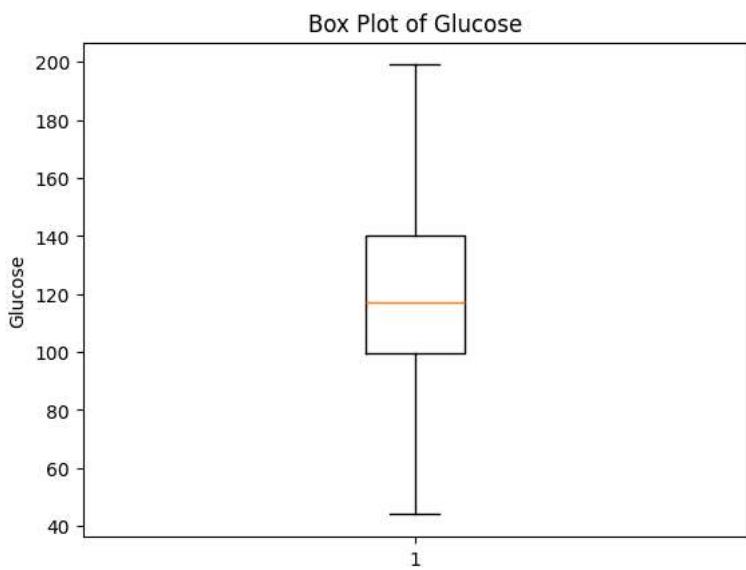
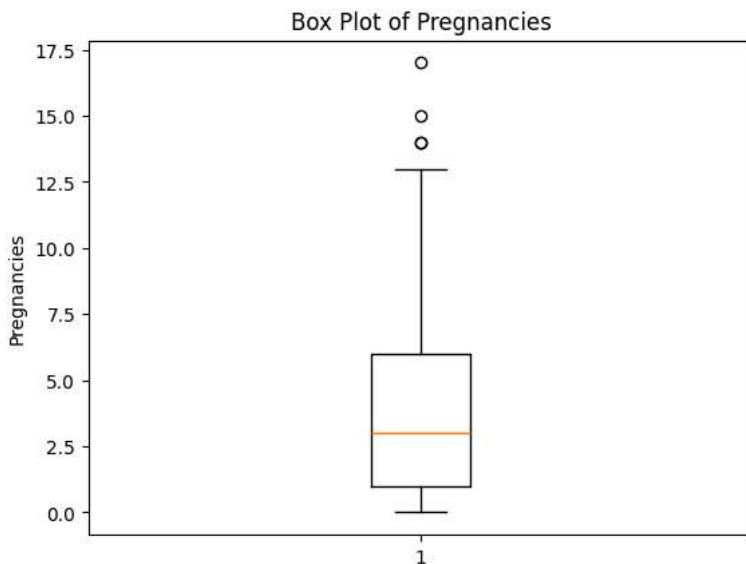
df.isnull().sum()
```

	0
Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

dtype: int64

```
numerical_cols = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']

for col in numerical_cols:
    plt.boxplot(df[col].dropna(), vert=True)
    plt.title(f"Box Plot of {col}")
    plt.ylabel(col)
    plt.grid(False)
    plt.show()
```

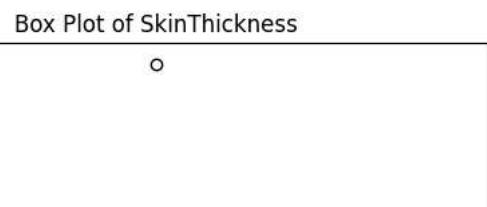
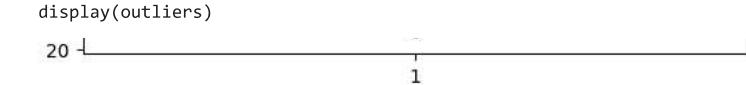


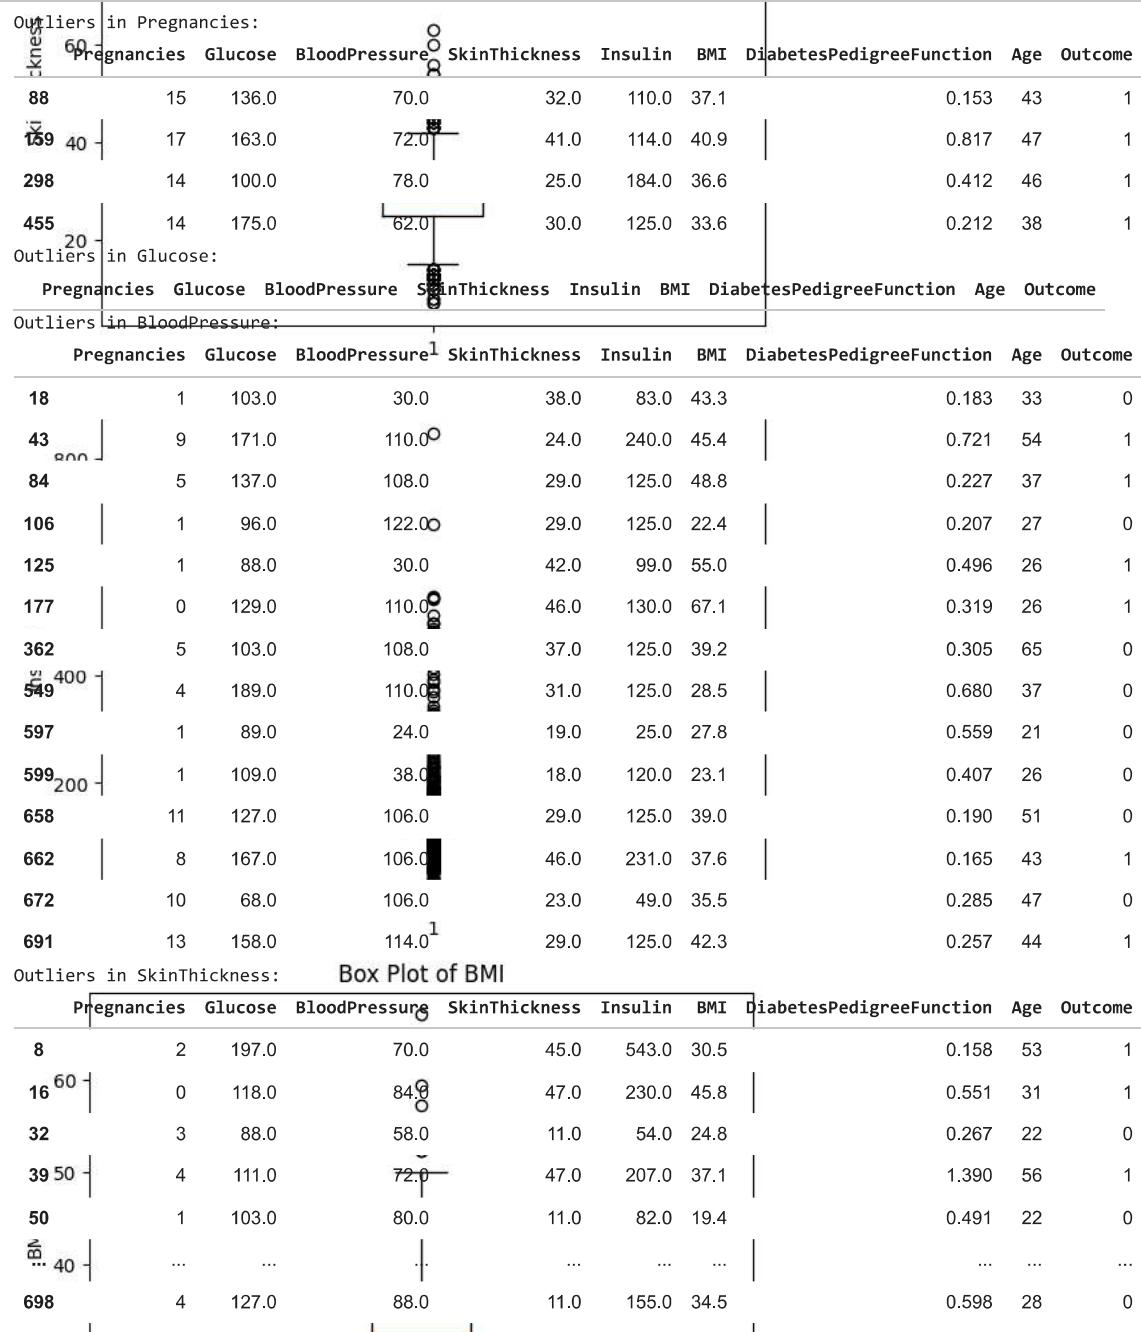
```
numerical_cols = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']

for col in numerical_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
    print(f"Outliers in {col}:")
    display(outliers)
```



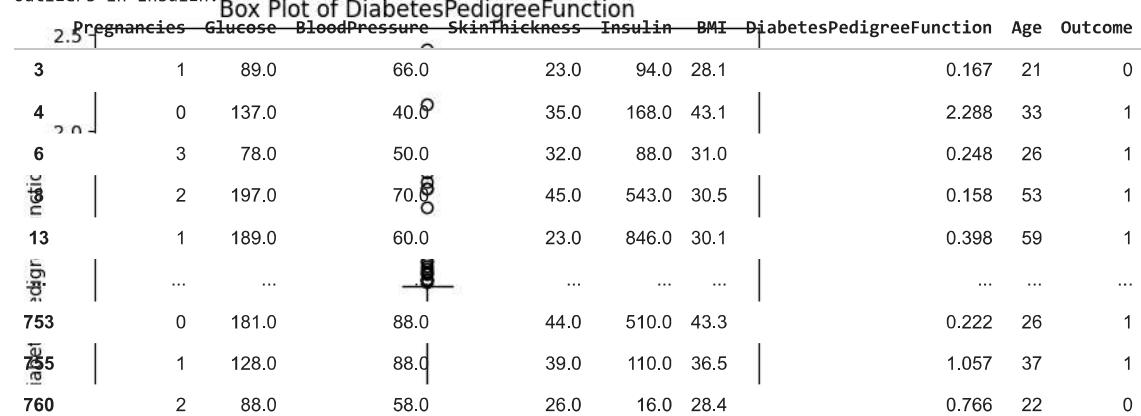


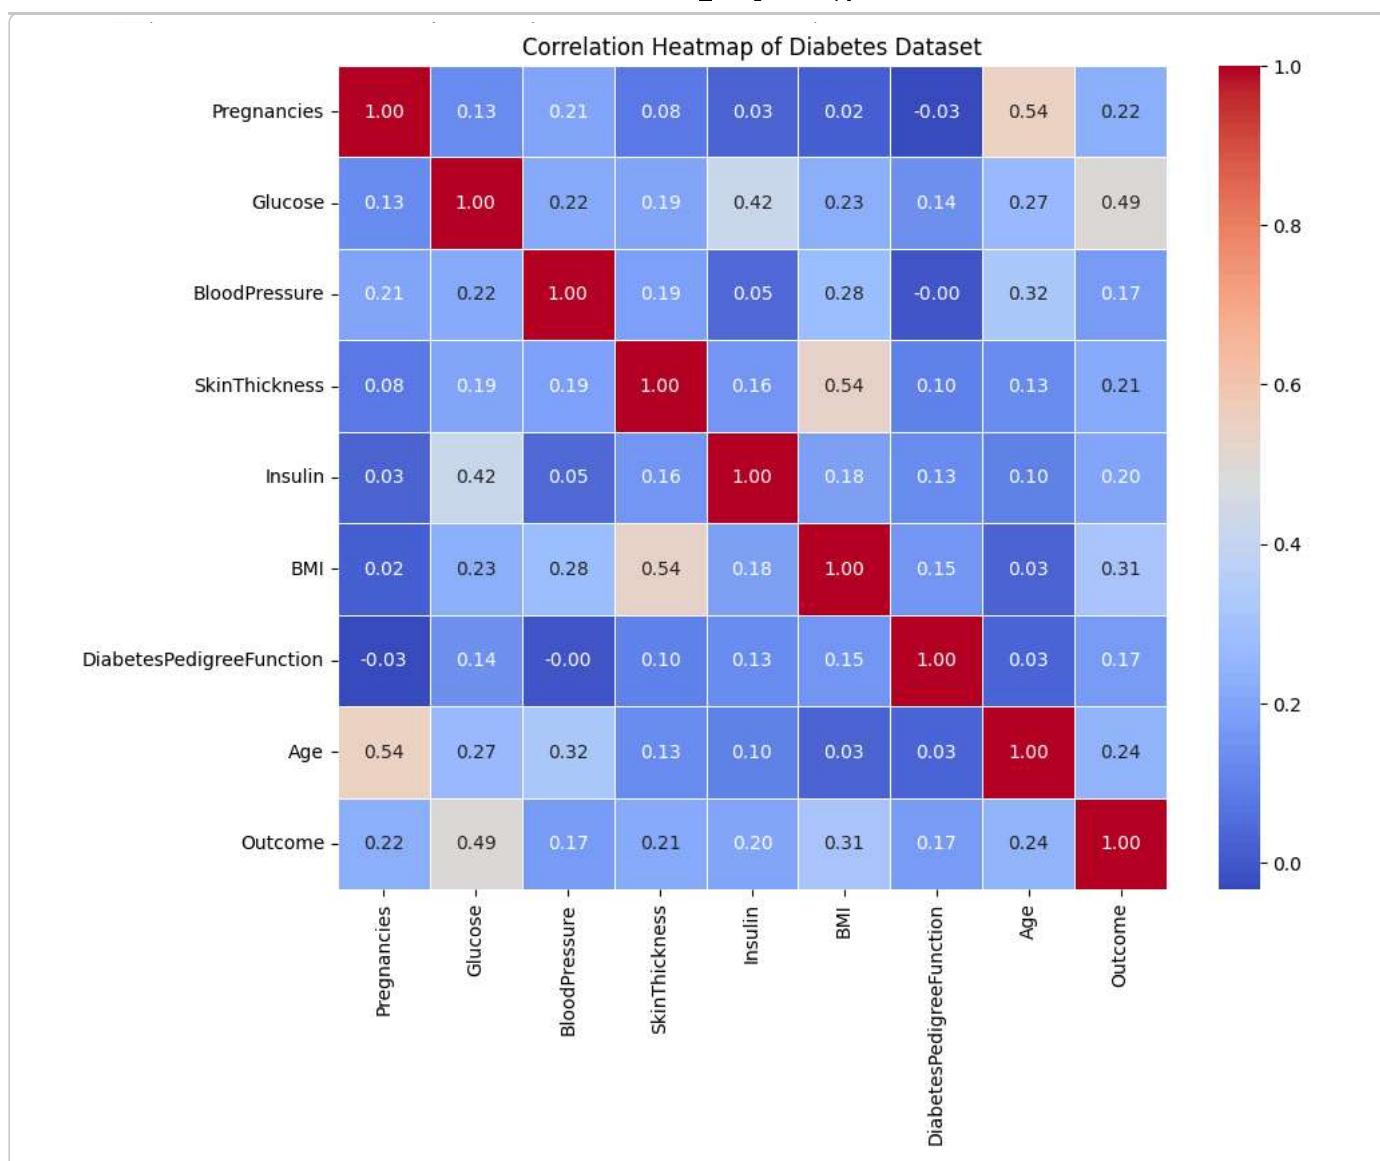
```
correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=.5)
plt.title('Correlation Heatmap of Diabetes Dataset')
plt.show()
```

87 rows × 9 columns

1

Outliers in Insulin:





292	2	128.0	78.0	37.0	182.0	43.300000		1.224	31	1		
<code>df.isnull().sum()</code>												
330	8	118.0	0	72.0	19.0	125.0	23.100000		1.476	46	0	
		Pregnancies	0		82.0	48.0	465.0	38.400000		2.137	25	1
371		Glucose	0	118.0	64.0	23.0	89.0	32.457464		1.731	21	0
		BloodPressure	0		62.0	18.0	59.0	25.100000		1.268	25	0
395		SkinThickness	2	127.0	58.0	24.0	275.0	27.700000		1.600	25	0
		Insulin	0		78.0	63.0	14.0	59.400000		2.420	25	1
534		BMI	1	77.0	56.0	30.0	56.0	33.300000		1.251	24	0
		DiabetesPedigreeFunction	0		52.0	22.0	115.0	28.500000		1.699	25	0
606		Age	1	181.0	78.0	42.0	293.0	40.000000		1.258	22	1
		Outcome	0		82.0	24.0	125.0	28.200000		1.282	50	1
621	2	92.0		76.0	20.0	125.0	24.200000		1.698	28	0	
<code>df.astype(int64)</code>												
661	1	199.0		76.0	43.0	125.0	42.900000		1.394	22	1	

```
df.to_csv('cleaned_dataset.csv', index=False) # Save the cleaned dataframe to a new CSV file
```

```
661      1    199.0        76.0       43.0     125.0    42.900000
from sklearn.linear_model import LogisticRegression      # For logistic regression model
from sklearn.tree import DecisionTreeClassifier        # For decision tree model
from sklearn.ensemble import RandomForestClassifier    # For random forest model
from sklearn.svm import SVC                            # For support vector machine
```

```

import matplotlib.pyplot as plt          # Import pyplot for plotting
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_auc_score      #

# Split dataframe into features (X) and label (y)
X = df.drop('Outcome', axis=1)          # Select all columns except 'Outcome' as predict
y = df['Outcome']                      # Select the 'Outcome' column as the target variable

# Logistic Regression
logreg = LogisticRegression(max_iter=200) # Create a Logistic Regression classifier with
logreg.fit(X_train, y_train)            # Train the logistic regression model using the training data

# Decision Tree
dtree = DecisionTreeClassifier(random_state=42) # Initialize a decision tree classifier with a random state
dtree.fit(X_train, y_train)              # Train the decision tree model using the training data

# Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42) # Initialize a random forest with 100 trees and a random state
rf.fit(X_train, y_train)                # Train the random forest model using the training data

# Support Vector Machine (SVM)
svm = SVC(probability=True, random_state=42) # Initialize a support vector machine classifier with probability=True and a random state
svm.fit(X_train, y_train)                 # Train the SVM model using the training data

def eval_metrics(y_test, y_pred, model_name): # This function prints accuracy, precision, recall, F1 score, and confusion matrix
    print(f"\n{model_name} Results:")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Precision:", precision_score(y_test, y_pred))
    print("Recall:", recall_score(y_test, y_pred))
    print("F1 Score:", f1_score(y_test, y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

from sklearn.neighbors import KNeighborsClassifier # Import the k-nearest neighbors classifier for this assignment

# Initialize and fit k-NN model (using k=5 neighbors)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Predict on test set and evaluate
y_pred_knn = knn.predict(X_test)
eval_metrics(y_test, y_pred_knn, "k-Nearest Neighbors")

k-Nearest Neighbors Results:
Accuracy: 0.6818181818181818
Precision: 0.5510204081632653
Recall: 0.5
F1 Score: 0.5242718446601942
Confusion Matrix:
[[78 22]
 [27 27]]

```

```
from sklearn.naive_bayes import GaussianNB # Import the Gaussian Naive Bayes classifier fo

# Initialize and fit Naive Bayes model
gnb = GaussianNB()
gnb.fit(X_train, y_train)

# Predict on test set and evaluate
y_pred_gnb = gnb.predict(X_test)
eval_metrics(y_test, y_pred_gnb, "Naive Bayes (GaussianNB)")
```

Naive Bayes (GaussianNB) Results:
Accuracy: 0.7012987012987013
Precision: 0.5666666666666667
Recall: 0.6296296296296297
F1 Score: 0.5964912280701754
Confusion Matrix:
[[74 26]
[20 34]]

```
from xgboost import XGBClassifier # Import XGBoost, a powerful and efficient grad

# Initialize and fit XGBoost model
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb.fit(X_train, y_train)

# Predict on test set and evaluate
y_pred_xgb = xgb.predict(X_test)
eval_metrics(y_test, y_pred_xgb, "XGBoost")
```

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [07:47:52] WARNING: /workspace/src/learner.cc:790: Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)

XGBoost Results:  

Accuracy: 0.7597402597402597  

Precision: 0.660377358490566  

Recall: 0.6481481481481481  

F1 Score: 0.6542056074766355  

Confusion Matrix:  

[[82 18]  

[19 35]]
```

```
y_pred_logreg = logreg.predict(X_test) # Variable that stores Logistic Regression pred
```

```
y_pred_dtree = dtree.predict(X_test) # Variable that stores Decision Tree prediction
```

```
y_pred_rf = rf.predict(X_test) # Variable that stores Random Forest prediction
```

```
y_pred_svm = svm.predict(X_test) # Variable that stores SVM predictions
```

```
eval_metrics(y_test, y_pred_logreg, "Logistic Regression") # Display and print evaluation results for Logi
```

Logistic Regression Results:
Accuracy: 0.7012987012987013
Precision: 0.5869565217391305
Recall: 0.5
F1 Score: 0.54
Confusion Matrix:
[[81 19]
[27 27]]

```
eval_metrics(y_test, y_pred_dtree, "Decision Tree") # Display and print evaluation results for Deci
```

Decision Tree Results:
Accuracy: 0.6688311688311688
Precision: 0.5333333333333333
Recall: 0.4444444444444444
F1 Score: 0.4848484848484848
Confusion Matrix:
[[79 21]]

[30 24]]

```
eval_metrics(y_test, y_pred_rf, "Random Forest") # Display and print evaluation results for Random Forest

Random Forest Results:
Accuracy: 0.7597402597402597
Precision: 0.6808510638297872
Recall: 0.5925925925925926
F1 Score: 0.6336633663366337
Confusion Matrix:
[[85 15]
 [22 32]]
```

```
eval_metrics(y_test, y_pred_svm, "SVM") # Display and print evaluation results for SVM

SVM Results:
Accuracy: 0.7207792207792207
Precision: 0.6486486486486487
Recall: 0.4444444444444444
F1 Score: 0.5274725274725275
Confusion Matrix:
[[87 13]
 [30 24]]
```

```
from sklearn.model_selection import cross_val_score

# Cross-Validation for Logistic Regression
logreg_cv = cross_val_score(logreg, X, y, cv=5, scoring='accuracy')
print("Logistic Regression Cross-Val Accuracy: {:.3f} (+/- {:.3f})".format(logreg_cv.mean(), logreg_cv.std()))

Logistic Regression Cross-Val Accuracy: 0.772 (+/- 0.020)
```

```
# Cross-Validation for Decision Tree
dtree_cv = cross_val_score(dtree, X, y, cv=5, scoring='accuracy')
print("Decision Tree Cross-Val Accuracy: {:.3f} (+/- {:.3f})".format(dtree_cv.mean(), dtree_cv.std()))

Decision Tree Cross-Val Accuracy: 0.720 (+/- 0.041)
```

```
# Cross-Validation for Random Forest
rf_cv = cross_val_score(rf, X, y, cv=5, scoring='accuracy')
print("Random Forest Cross-Val Accuracy: {:.3f} (+/- {:.3f})".format(rf_cv.mean(), rf_cv.std()))

Random Forest Cross-Val Accuracy: 0.760 (+/- 0.030)
```

```
# Cross-Validation for SVM
svm_cv = cross_val_score(svm, X, y, cv=5, scoring='accuracy')
print("SVM Cross-Val Accuracy: {:.3f} (+/- {:.3f})".format(svm_cv.mean(), svm_cv.std()))

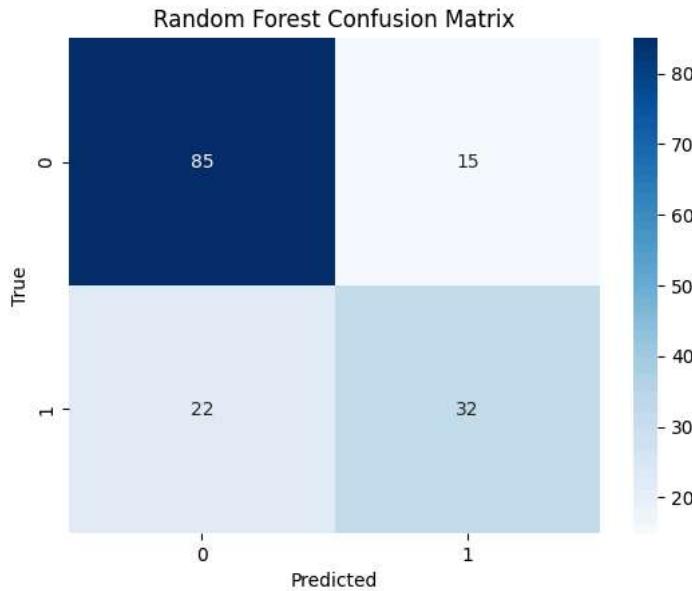
SVM Cross-Val Accuracy: 0.760 (+/- 0.023)
```

```
# Random Forest - Performance Metrics
rf_accuracy = accuracy_score(y_test, y_pred_rf)
rf_precision = precision_score(y_test, y_pred_rf)
rf_recall = recall_score(y_test, y_pred_rf)
rf_f1 = f1_score(y_test, y_pred_rf)
rf_auc = roc_auc_score(y_test, rf.predict_proba(X_test)[:, 1])
print("Random Forest:")
print("Accuracy:", rf_accuracy)
print("Precision:", rf_precision)
print("Recall:", rf_recall)
print("F1 Score:", rf_f1)
print("AUC:", rf_auc)

Random Forest:
Accuracy: 0.7597402597402597
Precision: 0.6808510638297872
Recall: 0.5925925925925926
F1 Score: 0.6336633663366337
AUC: 0.8200000000000001
```

```
import seaborn as sns

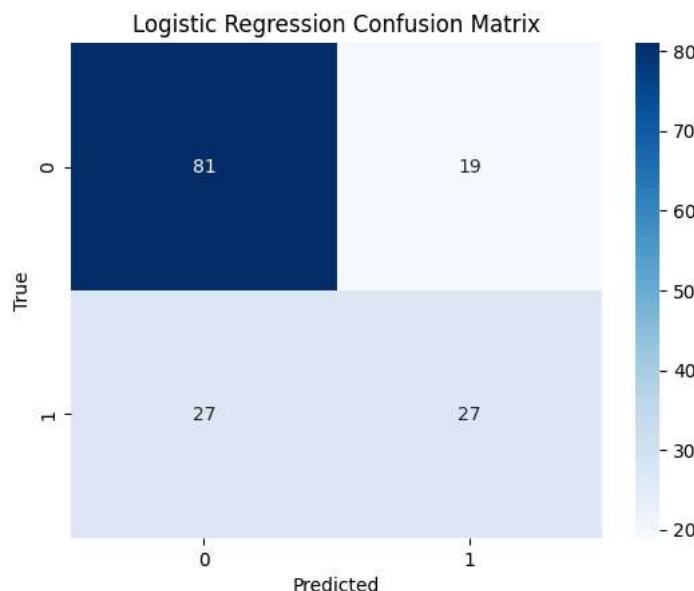
# Confusion matrix for Random Forest
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Blues')
plt.title('Random Forest Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



```
# Logistic Regression - Performance Metrics
logreg_accuracy = accuracy_score(y_test, y_pred_logreg)
logreg_precision = precision_score(y_test, y_pred_logreg)
logreg_recall = recall_score(y_test, y_pred_logreg)
logreg_f1 = f1_score(y_test, y_pred_logreg)
logreg_auc = roc_auc_score(y_test, logreg.predict_proba(X_test)[:, 1])
print("Logistic Regression:")
print("Accuracy:", logreg_accuracy)
print("Precision:", logreg_precision)
print("Recall:", logreg_recall)
print("F1 Score:", logreg_f1)
print("AUC:", logreg_auc)
```

Logistic Regression:
 Accuracy: 0.7012987012987013
 Precision: 0.5869565217391305
 Recall: 0.5
 F1 Score: 0.54
 AUC: 0.8127777777777777

```
# Confusion matrix for Logistic Regression
sns.heatmap(confusion_matrix(y_test, y_pred_logreg), annot=True, fmt='d', cmap='Blues')
plt.title('Logistic Regression Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

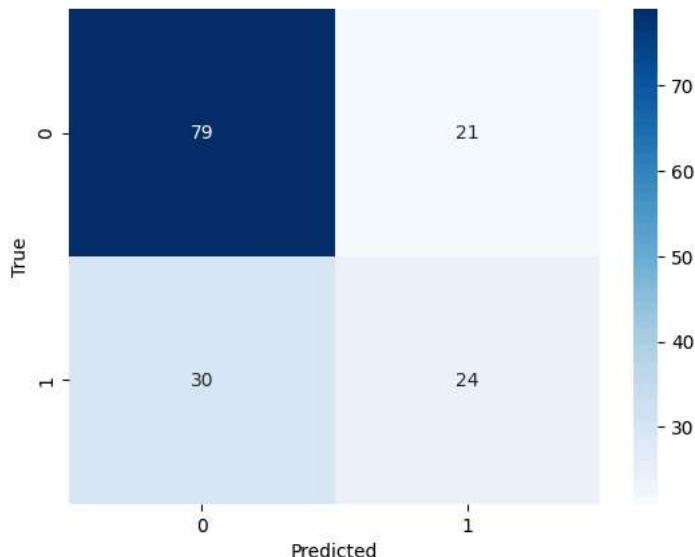


```
# Decision Tree - Performance Metrics
dtree_accuracy = accuracy_score(y_test, y_pred_dtree)
dtree_precision = precision_score(y_test, y_pred_dtree)
dtree_recall = recall_score(y_test, y_pred_dtree)
dtree_f1 = f1_score(y_test, y_pred_dtree)
dtree_auc = roc_auc_score(y_test, dtree.predict_proba(X_test)[:, 1])
print("Decision Tree:")
print("Accuracy:", dtree_accuracy)
print("Precision:", dtree_precision)
print("Recall:", dtree_recall)
print("F1 Score:", dtree_f1)
print("AUC:", dtree_auc)
```

```
Decision Tree:
Accuracy: 0.6688311688311688
Precision: 0.5333333333333333
Recall: 0.4444444444444444
F1 Score: 0.48484848484848486
AUC: 0.6172222222222222
```

```
# Confusion matrix for Decision Tree
sns.heatmap(confusion_matrix(y_test, y_pred_dtree), annot=True, fmt='d', cmap='Blues')
plt.title('Decision Tree Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

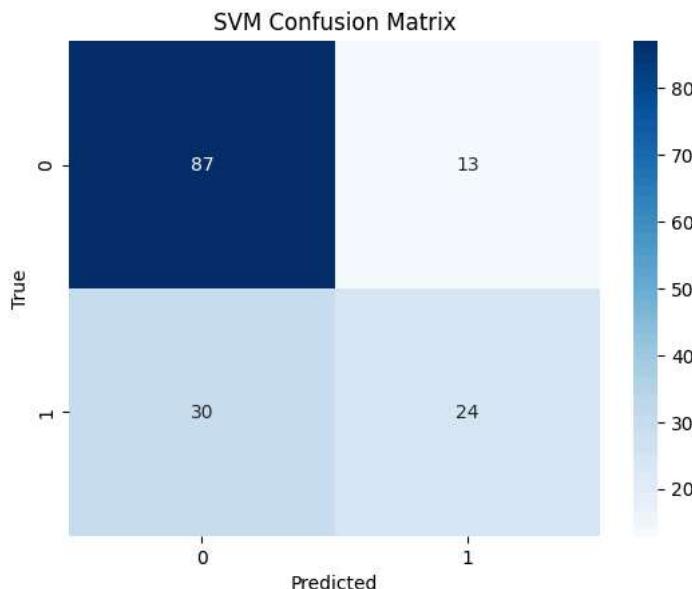
Decision Tree Confusion Matrix



```
# SVM - Performance Metrics
svm_accuracy = accuracy_score(y_test, y_pred_svm)
svm_precision = precision_score(y_test, y_pred_svm)
svm_recall = recall_score(y_test, y_pred_svm)
svm_f1 = f1_score(y_test, y_pred_svm)
svm_auc = roc_auc_score(y_test, svm.predict_proba(X_test)[:, 1])
print("Support Vector Machine (SVM):")
print("Accuracy:", svm_accuracy)
print("Precision:", svm_precision)
print("Recall:", svm_recall)
print("F1 Score:", svm_f1)
print("AUC:", svm_auc)
```

Support Vector Machine (SVM):
 Accuracy: 0.7207792207792207
 Precision: 0.6486486486486487
 Recall: 0.4444444444444444
 F1 Score: 0.5274725274725275
 AUC: 0.7811111111111111

```
# Confusion matrix for Support Vector Machine (SVM)
sns.heatmap(confusion_matrix(y_test, y_pred_svm), annot=True, fmt='d', cmap='Blues')
plt.title('SVM Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



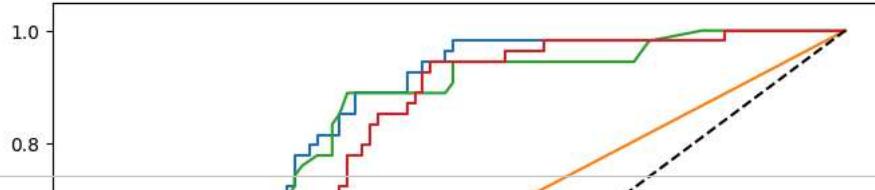
```
# Plot the ROC curves and AUC values for all models to visually compare their ability to distinguish between classes.
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

plt.figure(figsize=(8,6))
models = [
    (logreg, y_pred_logreg, "Logistic Regression"),
    (dtree, y_pred_dtree, "Decision Tree"),
    (rf, y_pred_rf, "Random Forest"),
    (svm, y_pred_svm, "SVM")
]

# Plot ROC curve for each model
for model, y_pred, name in models:
    # Get predicted probabilities for the positive class (1)
    if hasattr(model, "predict_proba"):
        y_score = model.predict_proba(X_test)[:,1]
    else:
        y_score = model.decision_function(X_test) # For SVM without probability
    fpr, tpr, _ = roc_curve(y_test, y_score)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

plt.plot([0, 1], [0, 1], 'k--') # Diagonal dashed line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for All Models')
plt.legend(loc='lower right')
plt.show()
```

ROC Curves for All Models



```
import pandas as pd

def predict_diabetes_from_input():
    print("Enter the following details:")

    pregnancies = int(input("Pregnancies: "))
    glucose = float(input("Glucose: "))
    bloodpressure = float(input("BloodPressure: "))
    skinthickness = float(input("SkinThickness: "))
    insulin = float(input("Insulin: "))
    bmi = float(input("BMI: "))
    dpf = float(input("DiabetesPedigreeFunction: "))
    age = int(input("Age: "))

    # Predict diabetes based on input features
    # (Implementation details omitted)
```