

AIM :- Implement K-Means Clustering Algorithm

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
df=pd.read_csv("/cars.csv")
```

```
df.shape
```

```
(38531, 30)
```

```
df.head()
```

	manufacturer_name	model_name	transmission	color	odometer_value	year_produced	engine_fuel	engine_has_gas	engine_type	...
0	Subaru	Outback	automatic	silver	190000	2010	gasoline	False	gasoline	
1	Subaru	Outback	automatic	blue	290000	2002	gasoline	False	gasoline	
2	Subaru	Forester	automatic	red	402000	2001	gasoline	False	gasoline	
3	Subaru	Impreza	mechanical	blue	10000	1999	gasoline	False	gasoline	
4	Subaru	Legacy	automatic	black	280000	2001	gasoline	False	gasoline	

5 rows × 30 columns

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38531 entries, 0 to 38530
Data columns (total 30 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   manufacturer_name 38531 non-null   object 
 1   model_name         38531 non-null   object 
 2   transmission       38531 non-null   object 
 3   color              38531 non-null   object 
 4   odometer_value     38531 non-null   int64  
 5   year_produced      38531 non-null   int64  
 6   engine_fuel        38531 non-null   object 
 7   engine_has_gas     38531 non-null   bool   
 8   engine_type        38531 non-null   object 
 9   engine_capacity    38521 non-null   float64
 10  body_type          38531 non-null   object 
 11  has_warranty       38531 non-null   bool   
 12  state              38531 non-null   object 
 13  drivetrain         38531 non-null   object 
 14  price_usd          38531 non-null   float64
 15  is_exchangeable    38531 non-null   bool   
 16  location_region    38531 non-null   object 
 17  number_of_photos   38531 non-null   int64  
 18  up_counter          38531 non-null   int64  
 19  feature_0           38531 non-null   bool   
 20  feature_1           38531 non-null   bool   
 21  feature_2           38531 non-null   bool   
 22  feature_3           38531 non-null   bool   
 23  feature_4           38531 non-null   bool   
 24  feature_5           38531 non-null   bool   
 25  feature_6           38531 non-null   bool   
 26  feature_7           38531 non-null   bool   
 27  feature_8           38531 non-null   bool   
 28  feature_9           38531 non-null   bool   
 29  duration_listed    38531 non-null   int64  
 dtypes: bool(13), float64(2), int64(5), object(10)
 memory usage: 5.5+ MB
```

```
df.describe()
```

	odometer_value	year_produced	engine_capacity	price_usd	number_of_photos	up_counter	duration_listed
count	38531.000000	38531.000000	38521.000000	38531.000000	38531.000000	38531.000000	38531.000000
mean	248864.638447	2002.943734	2.055161	6639.971021	9.649062	16.306091	80.577249
std	136072.376530	8.065731	0.671178	6428.152018	6.093217	43.286933	112.826569
min	0.000000	1942.000000	0.200000	1.000000	1.000000	1.000000	0.000000
25%	158000.000000	1998.000000	1.600000	2100.000000	5.000000	2.000000	23.000000
50%	250000.000000	2003.000000	2.000000	4800.000000	8.000000	5.000000	59.000000
75%	325000.000000	2009.000000	2.300000	8990.000000	12.000000	16.000000	91.000000
max	1000000.000000	2019.000000	8.000000	50000.000000	86.000000	1861.000000	2232.000000

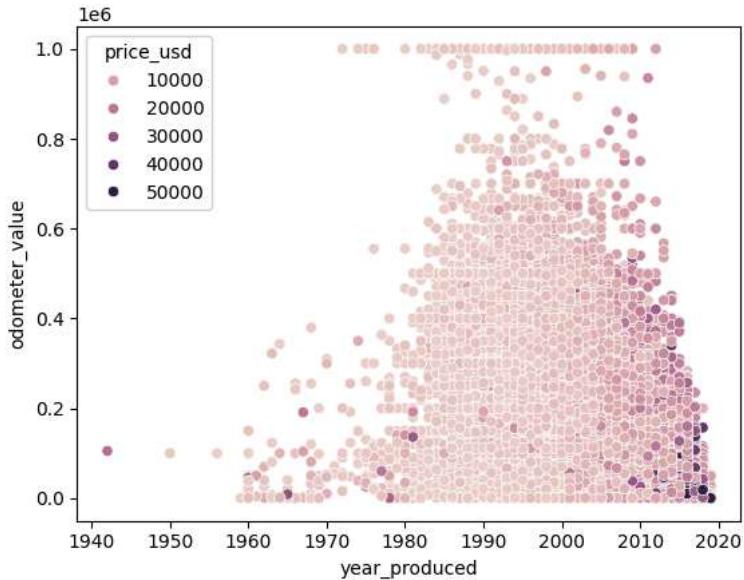
```
df.isnull().sum()
```

	0
manufacturer_name	0
model_name	0
transmission	0
color	0
odometer_value	0
year_produced	0
engine_fuel	0
engine_has_gas	0
engine_type	0
engine_capacity	10
body_type	0
has_warranty	0
state	0
drivetrain	0
price_usd	0
is_exchangeable	0
location_region	0
number_of_photos	0
up_counter	0
feature_0	0
feature_1	0
feature_2	0
feature_3	0
feature_4	0
feature_5	0
feature_6	0
feature_7	0
feature_8	0
feature_9	0
duration_listed	0

```
dtype: int64
```

```
sns.scatterplot(data=df,x='year_produced',y='odometer_value',hue='price_usd')
```

```
<Axes: xlabel='year_produced', ylabel='odometer_value'>
```



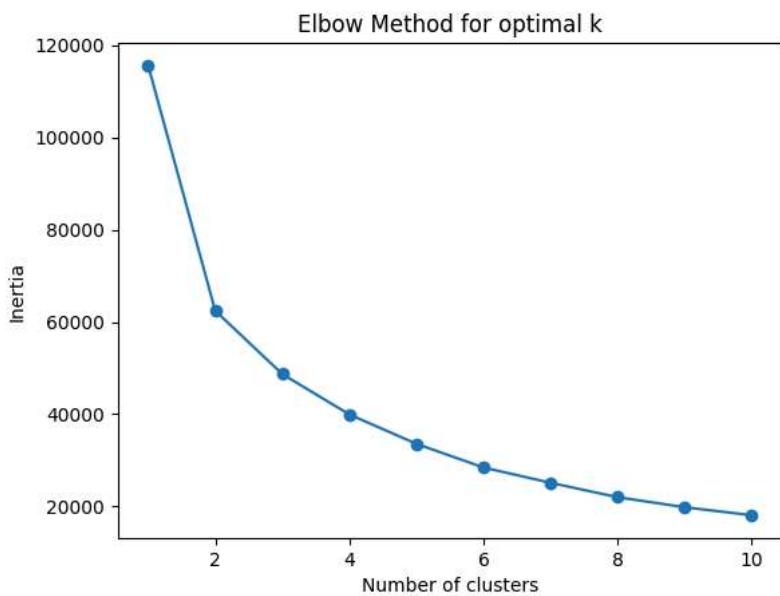
```
from sklearn.preprocessing import StandardScaler
#normalize the data
scaler=StandardScaler()
scaler.fit(df[['odometer_value','year_produced','price_usd']]) # intro to data for scalar
df_normalized=scaler.transform(df[['odometer_value','year_produced','price_usd']]) # makes it in one scale
print(df_normalized[:5])

[[-0.43260362  0.87485665  0.66272301]
 [ 0.30230894 -0.11700687 -0.25512656]
 [ 1.125411  -0.24098981 -0.59737555]
 [-1.75544621 -0.48895569  0.52255649]
 [ 0.22881768 -0.24098981 -0.70096654]]
```

```
from sklearn.cluster import KMeans
```

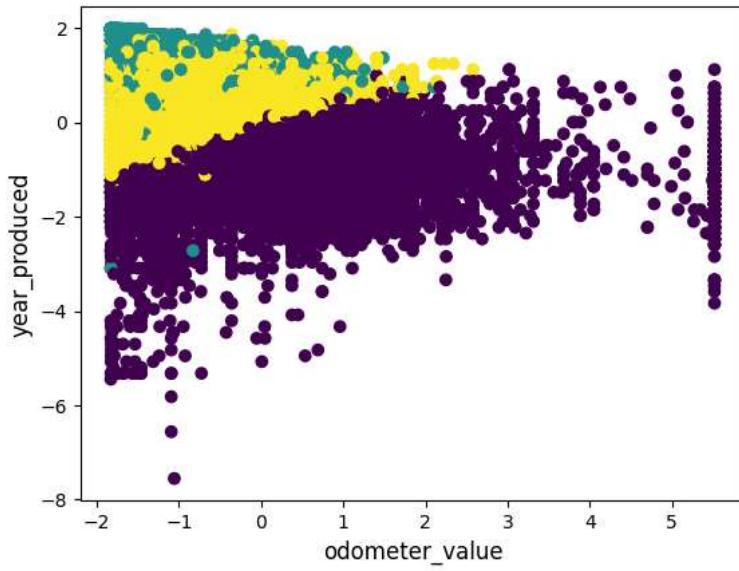
```
#calculate the within-clusters sum of square across different clusters counts
inertia=[]
for i in range(1,11):
    kmeans=KMeans(n_clusters=i,random_state=42,n_init='auto')
    kmeans.fit(df_normalized)
    inertia.append(kmeans.inertia_) # plot the elbow graph
```

```
plt.plot(range(1,11),inertia,marker='o')
plt.title('Elbow Method for optimal k')
plt.xlabel("Number of clusters")
plt.ylabel("Inertia")
plt.show()
```



```
#Elbow is at three clusters
kmeans=KMeans(n_clusters=3,random_state=42,n_init='auto')
clusters=kmeans.fit_predict(df_normalized)
```

```
#choose two dimensions to plot ( odometer_value and year_produced)
plt.scatter(df_normalized[:,0],df_normalized[:,1],c=clusters,cmap='viridis',marker='o')# 0- refers to odometer and 1- year
plt.xlabel('odometer_value',fontsize=12)
plt.ylabel('year_produced',fontsize=12)
plt.show()
```



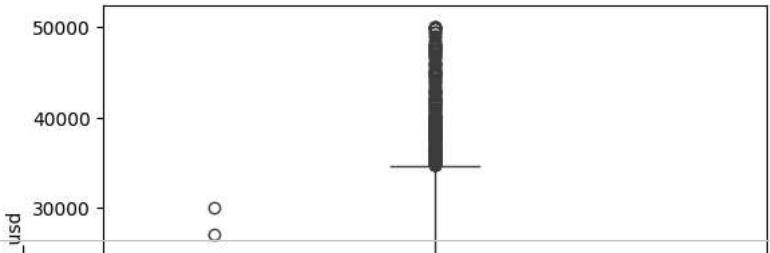
```
sns.boxplot(x=clusters,y=df['price_usd'])
summary_stats=df.groupby(clusters)['price_usd'].describe()
print(summary_stats)
```

```

      count      mean       std      min     25%     50%    75%  \
0  19276.0  2723.463427  1985.037525  1.00  1249.75  2300.0  3782.31
1   4568.0  19905.578398 7566.002147 9437.54 14746.28 17500.0 22753.48
2  14687.0   7654.290208 3088.330504  1.00  5500.00  7500.0  9700.00

      max
0  30000.0
1  50000.0
2  18500.0

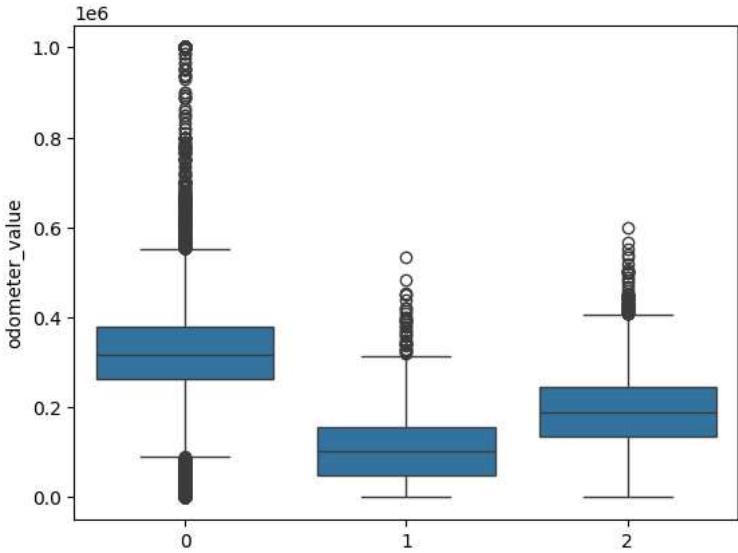
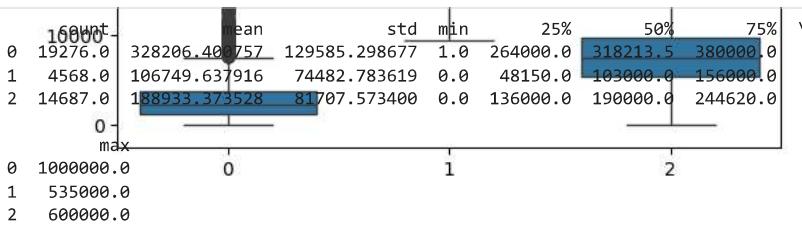
```



```

sns.boxplot(x=clusters,y=df['odometer_value'])
summary_stats=df.groupby(clusters)[['odometer_value']].describe()
print(summary_stats)

```

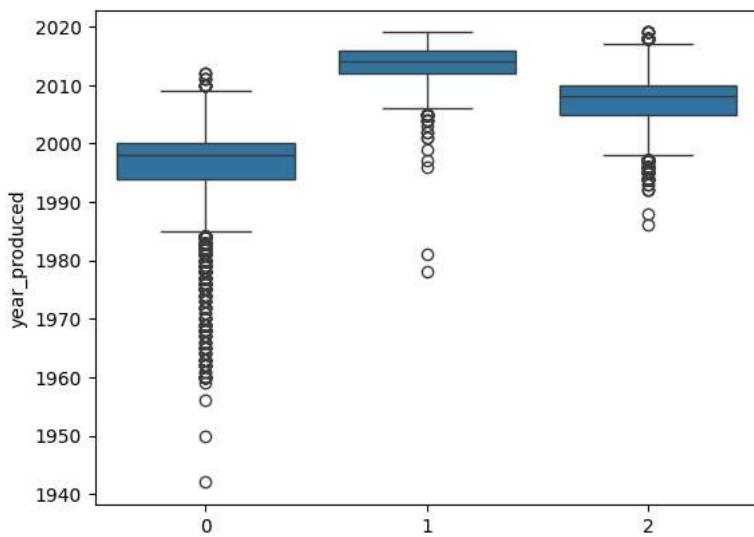


```

#CLUSTER ANALYSIS
sns.boxplot(x=clusters,y=df['year_produced'])
summary_stats=df.groupby(clusters)[['year_produced']].describe()
print(summary_stats)

```

	count	mean	std	min	25%	50%	75%	max
0	19276.0	1996.665335	5.411023	1942.0	1994.0	1998.0	2000.0	2012.0
1	4568.0	2014.069834	3.211429	1978.0	2012.0	2014.0	2016.0	2019.0
2	14687.0	2007.723361	3.994489	1986.0	2005.0	2008.0	2010.0	2019.0



```
#adding clusters column to the dataframe
df['cluster']=clusters
#now lets see the first five instances of the updated dataset
print(df.head(10))
```

0	Subaru	Outback	automatic	silver	190000			
1	Subaru	Outback	automatic	blue	290000			
2	Subaru	Forester	automatic	red	402000			
3	Subaru	Impreza	mechanical	blue	10000			
4	Subaru	Legacy	automatic	black	280000			
5	Subaru	Outback	automatic	silver	132449			
6	Subaru	Forester	automatic	black	318280			
7	Subaru	Legacy	automatic	silver	350000			
8	Subaru	Outback	automatic	grey	179000			
9	Subaru	Forester	automatic	silver	571317			
0	2010	gasoline	False	gasoline	2.5			
1	2002	gasoline	False	gasoline	3.0			
2	2001	gasoline	False	gasoline	2.5			
3	1999	gasoline	False	gasoline	3.0			
4	2001	gasoline	False	gasoline	2.5			
5	2011	gasoline	False	gasoline	2.5			
6	1998	gasoline	False	gasoline	2.5			
7	2004	gasoline	False	gasoline	2.5			
8	2010	gasoline	False	gasoline	2.5			
9	1999	gasoline	False	gasoline	2.5			
0	...	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	...
1	...	True	True	False	True	False	True	False
2	...	False	False	True	True	False	False	False
3	...	False	False	False	False	False	False	False
4	...	False	True	True	False	False	False	False
5	...	False	False	False	True	False	True	False
6	...	False	False	True	True	False	False	False
7	...	True	False	False	False	False	False	False
8	...	True	True	True	True	True	True	True
9	...	True	False	False	True	False	False	False
0	feature_8	feature_9	duration_listed	cluster				
1	True	True	16	2				
2	False	True	83	0				
3	True	True	151	0				
4	False	False	86	2				
5	False	True	7	0				
6	True	True	67	1				
7	True	True	307	0				
8	False	True	73	0				
9	True	True	87	2				
0	False	True	43	0				

[10 rows x 31 columns]

```
#load the original dataset(copy)
df_ori=pd.read_csv("/cars.csv")
df_ori.head() # without the clusters column
```

	manufacturer_name	model_name	transmission	color	odometer_value	year_produced	engine_fuel	engine_has_gas	engine_type	...
0	Subaru	Outback	automatic	silver	190000	2010	gasoline	False	gasoline	...
1	Subaru	Outback	automatic	blue	290000	2002	gasoline	False	gasoline	...
2	Subaru	Forester	automatic	red	402000	2001	gasoline	False	gasoline	...
3	Subaru	Impreza	mechanical	blue	10000	1999	gasoline	False	gasoline	...
4	Subaru	Legacy	automatic	black	280000	2001	gasoline	False	gasoline	...

5 rows × 30 columns

```
#Map the cluster number to a meaning
```

```
clusters_names_map={
    0:'Normal',
    1:'Recent',
    2:'Classic'
}
```

```
#create a new column with the cluster name
```

```
df_ori['cluster']=pd.Series(clusters).map(clusters_names_map)
```

```
#printing the new dataframe
```

```
df_ori.head(3)
```

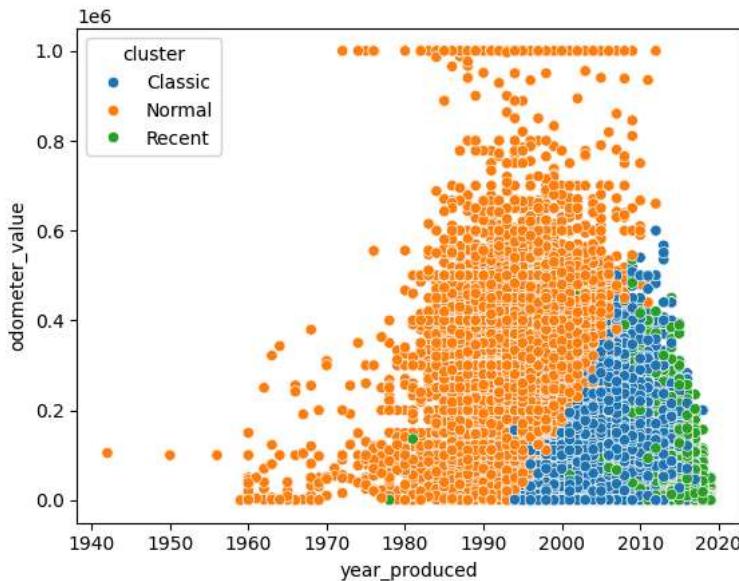
	manufacturer_name	model_name	transmission	color	odometer_value	year_produced	engine_fuel	engine_has_gas	engine_type	...
0	Subaru	Outback	automatic	silver	190000	2010	gasoline	False	gasoline	...
1	Subaru	Outback	automatic	blue	290000	2002	gasoline	False	gasoline	...
2	Subaru	Forester	automatic	red	402000	2001	gasoline	False	gasoline	...

3 rows × 31 columns

```
#plot tye new dataset we made with clusters
```

```
sns.scatterplot(data=df_ori,x='year_produced',y='odometer_value',hue='cluster')
```

```
<Axes: xlabel='year_produced', ylabel='odometer_value'>
```



```
fig=plt.figure(figsize=(12,9))
ax=fig.add_subplot(111,projection='3d')
```

```
#scatterplot using the first three features of the cars dataset
ax.scatter(df_normalized[:,0], #odometer_value
```

```
df_normalized[:,1], # year_produced  
df_normalized[:,2], # price_usd  
c=clusters, #use cluster label as color encoding  
cmap='viridis',  
marker='o')  
  
#set labels according to the features we used  
ax.set_xlabel('odometer_value')  
ax.set_ylabel('year_produced')  
ax.set_zlabel('price_usd')  
#title of the plot  
plt.title('3D Scatter plot of cars dataset')
```

3D Scatter plot of cars dataset

